

Chatbot System - Development Guide

Overview

We are building a chatbot with two main components:

1. **Data Ingestion System** (runs locally)
 2. **Chatbot** (deployed on Cloudflare)
-

1. Data Ingestion System

The purpose of this system is to prepare data (from PDFs) and store embeddings for chatbot retrieval.

Steps:

1. **Read PDF**
 - Extract text content from the PDF files.
 - Use libraries like [pdf2md](#).
 2. **Convert to Text & Clean**
 - Normalize formatting (remove extra spaces, newlines, etc.).
 3. **Chunking**
 - Split text into smaller chunks (e.g., 500–1000 characters with some overlap).
 - Helps embeddings and retrieval work better.
 4. **Create Embeddings**
 - Use **OpenAI Embeddings API** (e.g., `text-embedding-3-small` or `text-embedding-3-large`).
 5. **Store in Cloudflare Vectorize**
 - Push embeddings + metadata (document name, page number, chunk ID, original text) into **Cloudflare Vectorize**.
-

2. Chatbot (Cloudflare Worker)

The chatbot runs serverless using **Cloudflare Workers** and connects with external APIs for retrieval + inference.

Components:

1. **User Query Handling**
 - Accept user input (query).
 2. **Embedding User Query**
 - Create embedding for user query using **OpenAI Embeddings API**.
 3. **Vector Search**
 - Use **Cloudflare Vectorize** to search for most relevant text chunks.
 4. **Inference (Response Generation)**
 - Send query + retrieved chunks to **Groq API** (LLM inference).
 - Model generates the answer.
 5. **Return Response**
 - The chatbot returns the response to the user.
-

Tech Stack

- **Cloudflare Worker** → for chatbot runtime.
 - **Cloudflare Vectorize** → for vector database storage + retrieval.
 - **Groq API** → for fast LLM inference.
 - **OpenAI API** → for embeddings.
 - **Local NodeJS Script** → for data ingestion (PDF → Text → Embedding → Store).
-

Workflow Diagram (Simplified)

[PDF Files]



[Data Ingestion System (local)]

- Read PDF
- Chunk text
- Create embeddings (OpenAI)
- Store embeddings (Cloudflare Vectorize)

[Chatbot (Cloudflare Worker)]

- Take user query
- Create embedding (OpenAI)
- Search embeddings (CF Vectorize)
- Send context + query → Groq API
- Return answer